

# How To - AirFlow

## 1.1 What is Airflow?

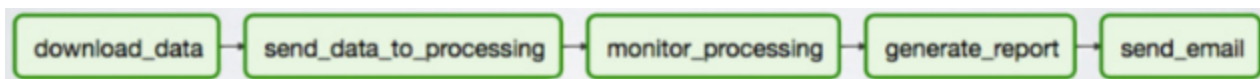
Airflow was developed by engineers at AirBnB to provide a standardized way to handle multiple ETL processes around an Enterprise Data Warehouse system. Airflow is written in Python but is language agnostic. It utilizes rabbitMQ, Jinja

- Airflow is a platform to programmatically author, schedule and monitor workflows or data pipelines.
- Use Airflow to author workflows as Directed Acyclic Graphs (DAGs) of tasks. The Airflow scheduler executes your tasks on an array of workers while following the specified dependencies.
- Easy to visualize pipelines running in production, monitor progress, and troubleshoot issues when needed.

## 1.2 What is a Workflow?

- a sequence of tasks
- started on a schedule or triggered by an event
- frequently used to handle big data processing pipelines

A typical workflows



- Download data from source
- send data somewhere else to process
- Monitor when the process is completed
- Get the result and generate the report
- Send the report out by email

## 1.3 How Apache Airflow Works

Apache Airflow accomplishes the tasks by taking DAG(Directed Acyclic Graphs) as an array of the workers, some of these workers have particularized contingencies.

- A workflow (data-pipeline) management system developed by Airbnb
- A framework to define tasks & dependencies in python
- Executing, scheduling, distributing tasks across worker nodes.
- View of present and past runs, logging feature
- Extensible through plugins
- Nice UI, possibility to define REST interface
- Interact well with database

## Principles

- **Dynamic:** Airflow pipelines are configuration as code (Python), allowing for dynamic pipeline generation. This allows for writing code that instantiates pipelines dynamically.
- **Extensible:** Easy to initiate the operators, executors and extend the library so that it fits the level of abstraction that suits your environment.
- **Elegant:** Airflow pipelines are lean and explicit. Parameterizing your scripts is built into the core of Airflow using the powerful **Jinja** templating engine.
- **Scalable:** The architecture of Airflow composed of standardized units which also use messaging technique for queuing the number of workers and moreover it is scalable to infinity.

## Use cases

Airflow as an orchestration tool to coordinate work done by other services. It's not a data streaming solution—even though tasks can exchange some metadata, they do not move data among themselves. Example types of use cases suitable for Airflow:

- ETL (extract, transform, load) jobs - extracting data from multiple sources, transforming for analysis and loading it into a data store

- Machine Learning pipelines
- Data warehousing
- Orchestrating automated testing
- Performing backups

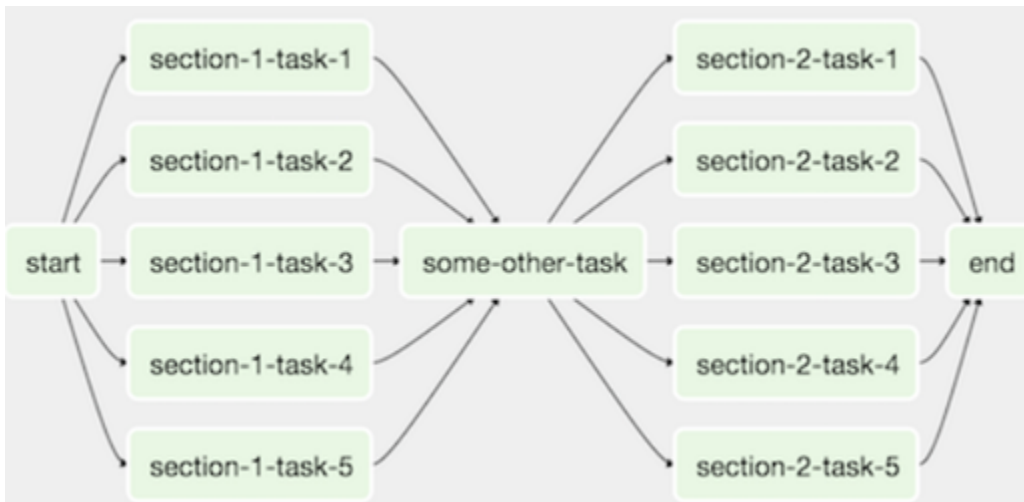
It is generally best suited for regular operations which can be scheduled to run at specific times.

## 2. Core concepts

### 2.1 Airflow DAG

First of them is the DAG - short for Directed Acyclic Graph. It's a collection of all the tasks you want to run, taking into account dependencies between them.

- Workflow as a Directed Acyclic Graph (DAG) with multiple tasks which can be executed independently.
- Airflow DAGs are composed of Tasks.



### 2.2 Airflow operators

While DAGs describe how to run a workflow, Airflow operators determine what actually gets done. There are several types of operators:

**OPERATOR** : Action operators which perform a single operation and return e.g. `BashOperator`, `GoogleCloudStorageDownloadOperator`

**Sensor** : sensors which pause the execution until certain criteria are met, such as until a certain key appears in S3 e.g. `GoogleCloudStorageObjectSensor`

**Executor** : transfer operators which connect 2 services and enable sending data between them e.g. `GoogleCloudStorageToS3Operator`

An operator is simply a Python class with an “execute()” method, which gets called when it is being run.

```

class ExampleOperator(BaseOperator):
    def execute(self, context):
        # Do something
        pass
  
```

In the same vein a sensor operator is a Python class with a “poke()” method, which gets called repeatedly until “True” is returned.

```
class ExampleSensorOperator(BaseSensorOperator):
    def poke(self, context):
        # Check if the condition occurred
        return True
```

## 2.3 Task

In order to execute an operator we need to create a task, which is a representation of the operator with a particular set of input arguments.

example: we have one BashOperator, but we can create three different “bash tasks” in a DAG, where each task is passed a different bash command to execute. So far we have the DAG, operators and tasks.

### Source

<https://airflow.apache.org/>

<https://www.xenonstack.com/insights/apache-airflow/>

<https://www.applydatascience.com/airflow/airflow-tutorial-introduction/>

<https://www.polidea.com/blog/apache-airflow-tutorial-and-beginners-guide/>